

Introduction to Theano

Frédéric Bastien

(slides highly copied from previous tutorial by Ian G.)

September 25, 2014

High level

- ▶ Overview of library (3 min)
- ▶ Building expressions (30 min)
- ▶ Compiling and running expressions (30 min)
- ▶ Modifying expressions (25 min)
- ▶ Debugging (30 min)
- ▶ Citing Theano (2 min)

Overview of Library

Theano is many things

- ▶ Language
- ▶ Compiler
- ▶ Python library

Overview

Theano language:

- ▶ Operations on scalar, vector, matrix, tensor, and sparse variables
- ▶ Linear algebra
- ▶ Element-wise nonlinearities
- ▶ Convolution
- ▶ Extensible

Overview

Using Theano:

- ▶ define expression $f(x, y) = x + y$
- ▶ compile expression

```
int f(int x, int y){  
    return x + y;  
}
```

- ▶ execute expression

```
>>> f(1, 2)  
3
```

Building expressions

- ▶ Scalars
- ▶ Vectors
- ▶ Matrices
- ▶ Tensors
- ▶ Reduction
- ▶ Dimshuffle

Scalar math

Using Theano:

- ▶ define expression $f(x, y) = x + y$
- ▶ compile expression

```
from theano import tensor as T
x = T.scalar()
y = T.scalar()
z = x+y
w = z*x
a = T.sqrt(w)
b = T.exp(a)
c = a ** b
d = T.log(c)
```

Vector math

```
from theano import tensor as T
x = T.vector()
y = T.vector()
# Scalar math applied elementwise
a = x * y
# Vector dot product
b = T.dot(x, y)
# Broadcasting
c = a + b
```


Matrix math

```
from theano import tensor as T
x = T.matrix()
y = T.matrix()
a = T.vector()
# Matrix-matrix product
b = T.dot(x, y)
# Matrix-vector product
c = T.dot(x, a)
```

Tensors

Using Theano:

- ▶ define expression $f(x, y) = x + y$
- ▶ compile expression
 - ▶ Dimensionality defined by length of “broadcastable” argument
 - ▶ Can add (or do other elemwise op) on two tensors with same dimensionality
 - ▶ Duplicate tensors along broadcastable axes to make size match

```
from theano import tensor as T
tensor3 = T.TensorType(
    broadcastable=(False, False, False),
    dtype='float32')
x = tensor3()
```

Reductions

Using Theano:

- ▶ define expression $f(x, y) = x + y$
- ▶ compile expression

```
from theano import tensor as T
tensor3 = T.TensorType(
    broadcastable=(False, False, False),
    dtype='float32')
x = tensor3()
total = x.sum()
marginals = x.sum(axis=(0, 2))
mx = x.max(axis=1)
```

Dimshuffle

```
from theano import tensor as T
tensor3 = T.TensorType(broadcastable=(False, False,
x = tensor3()
y = x.dimshuffle((2, 1, 0))
a = T.matrix()
b = a.T
# Same as b
c = a.dimshuffle((0, 1))
# Adding to larger tensor
d = a.dimshuffle((0, 1, 'x'))
e = a + d
```

Exercices

Work through the “01_buildbing_expressions” directory now.

Available at

“git https://github.com/nouiz/ccw_tutorial_theano.git”.

Compiling and running expression

- ▶ `theano.function`
- ▶ shared variables and updates
- ▶ compilation modes
- ▶ compilation for GPU
- ▶ optimizations

theano.function

```
>>> from theano import tensor as T
>>> x = T.scalar()
>>> y = T.scalar()
>>> from theano import function
>>> # first arg is list of SYMBOLIC inputs
>>> # second arg is SYMBOLIC output
>>> f = function([x, y], x + y)
>>> # Call it with NUMERICAL values
>>> # Get a NUMERICAL output
>>> f(1., 2.)
array(3.0)
```

Shared variables

- ▶ It's hard to do much with purely functional programming
- ▶ “shared variables” add just a little bit of imperative programming
- ▶ A “shared variable” is a buffer that stores a numerical value for a Theano variable
- ▶ Can write to as many shared variables as you want, once each, at the end of the function
- ▶ Modify outside Theano function with `get_value()` and `set_value()` methods.

Shared variable example

```
>>> from theano import shared
>>> x = shared(0.)
>>> from theano.compat.python2x import OrderedDict
>>> updates = OrderedDict()
>>> updates[x] = x + 1
>>> f = function([], updates=updates)
>>> f()
>>> x.get_value()
1.0
>>> x.set_value(100.)
>>> f()
>>> x.get_value()
101.0
```

Which dict?

- ▶ Use `theano.compat.python2x.OrderedDict`
- ▶ Not `collections.OrderedDict`
 - ▶ This isn't available in older versions of python, and will limit the portability of your code
- ▶ Not `{}` aka dict
 - ▶ The iteration order of this built-in class is not deterministic (thanks, Python!) so if Theano accepted this, the same script could compile different C programs each time you run it

Compilation modes

- ▶ Can compile in different modes to get different kinds of programs
- ▶ Can specify these modes very precisely with arguments to `theano.function`
- ▶ Can use a few quick presets with environment variable flags

Example preset compilation modes

- ▶ `FAST_RUN`: default. Spends a lot of time on compilation to get an executable that runs fast.
- ▶ `FAST_COMPILE`: Doesn't spend much time compiling. Executable usually uses python instead of compiled C code. Runs slow.
- ▶ `DEBUG_MODE`: Adds lots of checks. Raises error messages in situations other modes regard as fine.

Compilation for GPU

- ▶ Theano current back-end only supports 32 bit on GPU
- ▶ CUDA supports 64 bit, but is slow in gamer card
- ▶ T.fscalar, T.fvector, T.fmatrix are all 32 bit
- ▶ T.scalar, T.vector, T.matrix resolve to 32 bit or 64 bit depending on theano's floatX flag
- ▶ floatX is float64 by default, set it to float32
- ▶ Set device flag to gpu (or a specific gpu, like gpu0)

Optimizations

- ▶ Theano changes the symbolic expressions you write before converting them to C code
- ▶ It makes them faster
 - ▶ $(x+y)+(x+y) \rightarrow 2(x+y)$
- ▶ It makes them more stable
 - ▶ $\exp(a)/\exp(a).sum() \rightarrow \text{softmax}(a)$

Optimizations

- Sometimes optimizations discard error checking and produce incorrect output rather than an exception

```
>>> x = T.scalar()
>>> f = function([x], x/x)
>>> f(0.)
array(1.0)
```

Exercises

Work through the “02_compiling_and_running” directory now

Modifying expressions

- ▶ The grad method
- ▶ Variable nodes
- ▶ Types
- ▶ Ops
- ▶ Apply nodes

The grad method

```
>>> x = T.scalar('x')
>>> y = 2. * x
>>> g = T.grad(y, x)
>>> from theano.printing import min_informative_str
>>> print min_informative_str(g)
```

- A. Elemwise{mul}
- B. Elemwise{second,no_inplace}
- C. Elemwise{mul,no_inplace}
- D. TensorConstant{2.0}
- E. x
- F. TensorConstant{1.0}

<D>

Theano Variables

- ▶ A Variable is a theano expression
- ▶ Can come from T.scalar, T.matrix, etc.
- ▶ Can come from doing operations on other Variables
- ▶ Every Variable has a type field, identifying its Type
e.g. `TensorType((True, False), 'float32')`
- ▶ Variables can be thought of as nodes in a graph

Ops

- ▶ An Op is any class that describes a mathematical function of some variables
- ▶ Can call the op on some variables to get a new variable or variables
- ▶ An Op class can supply other forms of information about the function, such as its derivatives

Apply nodes

- ▶ The Apply class is a specific instance of an application of an Op
- ▶ Notable fields:
 - ▶ op: The Op to be applied
 - ▶ inputs: The Variables to be used as input
 - ▶ outputs: The Variables produced
- ▶ Variable.owner identifies the Apply that created the variable
- ▶ Variable and Apply instances are nodes and owner/inputs/outputs identify edges in a Theano graph

Exercises

Work through the “03_modifying” directory now

Debugging

- ▶ `DEBUG_MODE`
- ▶ Error message
- ▶ `theano.printing.debugprint`
- ▶ `min_informative_str`
- ▶ `compute_test_value`
- ▶ Accessing the `FunctionGraph`

Error message: code

```
import numpy as np
import theano
import theano.tensor as T
x = T.vector()
y = T.vector()
z = x + x
z = z + y
f = theano.function([x, y], z)
f(np.ones((2,)), np.ones((3,)))
```


Error message: 1st part

```
Traceback (most recent call last):
```

```
[...]
```

```
ValueError: Input dimension mis-match.
```

```
    (input[0].shape[0] = 3, input[1].shape[0] = 2)
```

```
Apply node that caused the error:
```

```
    Elemwise{add,no_inplace}(<TensorType(float64, vector),  
                             <TensorType(float64, vector),  
                             <TensorType(float64, vector)
```

```
Inputs types: [TensorType(float64, vector),  
               TensorType(float64, vector),  
               TensorType(float64, vector)]
```

```
Inputs shapes: [(3,), (2,), (2,)]
```

```
Inputs strides: [(8,), (8,), (8,)]
```

```
Inputs scalar values: ['not scalar', 'not scalar']
```

Error message: 2st part

Debugprint of the apply node:

```
Elemwise{add,no_inplace} [@A] <TensorType(float64 , vector)>  
|<TensorType(float64 , vector)> [@B] <TensorType(float64 , vector)>  
|<TensorType(float64 , vector)> [@C] <TensorType(float64 , vector)>  
|<TensorType(float64 , vector)> [@C] <TensorType(float64 , vector)>
```

HINT: Re-running with most Theano optimization disabled could give you a back-traces when this node was created. This can be done with by setting the Theano flags `optimizer=fast_compile`

Error message: optimizer=fast_compile

Backtrace when the node is created:

File `"test.py"`, line 7, in `<module>`

`z = z + y`

File `"/home/nouiz/src/Theano/theano/tensor/var.py"`

`return theano.tensor.basic.add(self, other)`

debugprint

```
>>> from theano.printing import debugprint
>>> debugprint(a)
Elemwise{mul,no_inplace} [@A] ' '
| TensorConstant{2.0} [@B]
| Elemwise{add,no_inplace} [@C] 'z'
|<TensorType(float64, scalar)> [@D]
|<TensorType(float64, scalar)> [@E]
```

min_informative_str

```
>>> x = T.scalar()
>>> y = T.scalar()
>>> z = x + y
>>> z.name = 'z'
>>> a = 2. * z
>>> from theano.printing import min_informative_str
>>> print min_informative_str(a)
A. Elemwise{mul,no_inplace}
B. TensorConstant{2.0}
C. z
```

compute_test_value

```
>>> from theano import config
>>> config.compute_test_value = 'raise'
>>> x = T.vector()
>>> import numpy as np
>>> x.tag.test_value = np.ones((2,))
>>> y = T.vector()
>>> y.tag.test_value = np.ones((3,))
>>> x + y
...
ValueError: Input dimension mismatch.
(input[0].shape[0] = 2, input[1].shape[0] = 3)
```

Accessing a function's fgraph

```
>>> x = T.scalar()
>>> y = x / x
>>> f = function([x], y)
>>> debugprint(f.maker.fgraph.outputs[0])
DeepCopyOp  [@A]  ' '
| TensorConstant{1.0}  [@B]
```

Exercises

Work through the “04_debugging” directory now

Citing Theano

- ▶ Please cite both of the following papers in all work that uses Theano:
- ▶ Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Bergstra, James, Goodfellow, Ian, Bergeron, Arnaud, Bouchard, Nicolas, and Bengio, Yoshua. Theano: new features and speed improvements. Deep Learning and Unsupervised Feature Learning NIPS 2012 Workshop, 2012.
- ▶ Bergstra, James, Breuleux, Olivier, Bastien, Frédéric, Lamblin, Pascal, Pascanu, Razvan, Desjardins, Guillaume, Turian, Joseph, Warde-Farley, David, and Bengio, Yoshua. Theano: a CPU and GPU math expression compiler. In Proceedings of the Python for Scientific Computing Conference (SciPy), June 2010. Oral Presentation.

Example acknowledgments

We would like to thank the developers of Theano
`citep{bergstra+al:2010-scipy,Bastien-Theano-2012}`, Pylearn2
`citep{pylearn2_arxiv_2013}`. We would also like to thank NSERC,
Compute Canada, and Calcul Québec for providing computational
resources.

Questions?